

Pilot Evidence Packet v1.0

SHADOW Mode External Log Ingestion

MathLedger Project

2025-12-17

Status: COMPLETE / FROZEN (SHADOW MODE)

1 Purpose

This document describes a **SHADOW ingest interface for external logs**. The pilot ingestion adapter receives external JSON/JSONL logs, attaches toolchain provenance, and routes data to an isolated namespace (`governance.external.pilot.*`).

Mode: SHADOW (observational only)

Action: LOGGED_ONLY

Enforcement: None

2 What This Demonstrates

| Property | Verification |
|---------------------|---|
| Namespace isolation | External data routed to <code>governance.external.pilot.*</code> only |
| Non-interference | Internal metrics, verifier state, CAL-EXP surfaces unchanged |
| Determinism | Repeated runs produce byte-identical outputs (timestamp-stripped) |
| Schema separation | Pilot schema does not extend v0.1 contract schemas |

3 What This Does NOT Demonstrate

| Non-Claim | Explanation |
|-------------|---|
| Validity | Pilot outputs do not constitute experimental evidence |
| Correctness | Execution does not imply correctness |
| Learning | No system learning or adaptation claims |
| Gating | Pilot outputs never block execution |
| Readiness | Does not demonstrate production readiness |

4 Smoke Proof (Copy/Paste)

Run from repository root:

```
uv run python -c "
import json, tempfile
from pathlib import Path
from backend.health.pilot_external_ingest_adapter import (
```

```

    ingest_external_log, wrap_for_evidence_pack,
    attach_to_manifest, copy_to_evidence_pack,
)
with tempfile.TemporaryDirectory() as t:
    tmp = Path(t)
    log_json = tmp / 'sample_external.json'
    log_json.write_text(json.dumps({
        'log_type': 'pilot_demo',
        'entries': [{'event': 'smoke', 'value': 1}]
    }))
    pack = tmp / 'evidence_pack'
    pack.mkdir()
    r = ingest_external_log(log_json)
    e = wrap_for_evidence_pack(r, log_json)
    m = attach_to_manifest({'governance': {}}, [e])
    print(f'result={r["result"]} mode={m["governance"]["external_pilot"] ["mode"]}')
"
```

Expected output:

```
result=SUCCESS mode=SHADOW
```

5 Evidence: Test Suite

5.1 Isolation Tests

File: tests/integration/test_external_pilot_signal_isolation.py

- 11 test classes, 38+ test methods
- Covers: valid, malformed, empty, adversarial, oversized logs
- Guarantees: namespace isolation, verifier invariance, input purity
- Meta-tests: registry tripwire (minimum 35 tests threshold)

5.2 Text Neutrality Tests

File: tests/policy/test_pilot_text_neutrality.py

- 3 test classes, 8 test methods
- Guarantees: no banned alarm words, no evaluative terms
- Structured tokens (PASS/WARN) explicitly allowed

5.3 Summary

| Test Suite | Classes | Methods |
|------------------|-----------|------------|
| Signal Isolation | 11 | 38+ |
| Text Neutrality | 3 | 8 |
| Total | 14 | 46+ |

6 Language Posture

All pilot components use neutral language:

- Warning strings avoid alarm words
- Structured result codes (SUCCESS, SCHEMA_INVALID, etc.)
- Disclaimer included in all outputs

Pilot provenance disclaimer (required in all artifacts):

“PILOT PROVENANCE: Binds artifact to toolchain state at ingestion. NOT experiment provenance. Does not imply experimental validity, reproducibility, or parity with CAL-EXP runs.”

7 Contact / Next Step

Safe Pilot Handshake: External partners may provide JSON/JSONL logs for ingestion via the pilot adapter. All ingested data is isolated to `governance.external_pilot.*` and cannot affect internal metrics, verifier behavior, or CAL-EXP verification surfaces. To initiate a pilot handshake, contact the MathLedger governance council with sample log structure and expected schema.

Document: PILOT_EVIDENCE_PACKET v1.0

Mode: SHADOW (observational only)

Authority: PILOT_NON_INTERFERENCE_ATTESTATION.md

Generated: 2025-12-17